

佟达



TENSORFLOW实现卷积神经网络

2016.10.28

回顾 - 神经网络

详细请看《程序员眼中的深度学习》

多层感知机——任务定义



$$L_{\text{hidden}} = \text{sigmoid}(x \cdot W_{\text{hidden}} + b_{\text{hidden}})$$

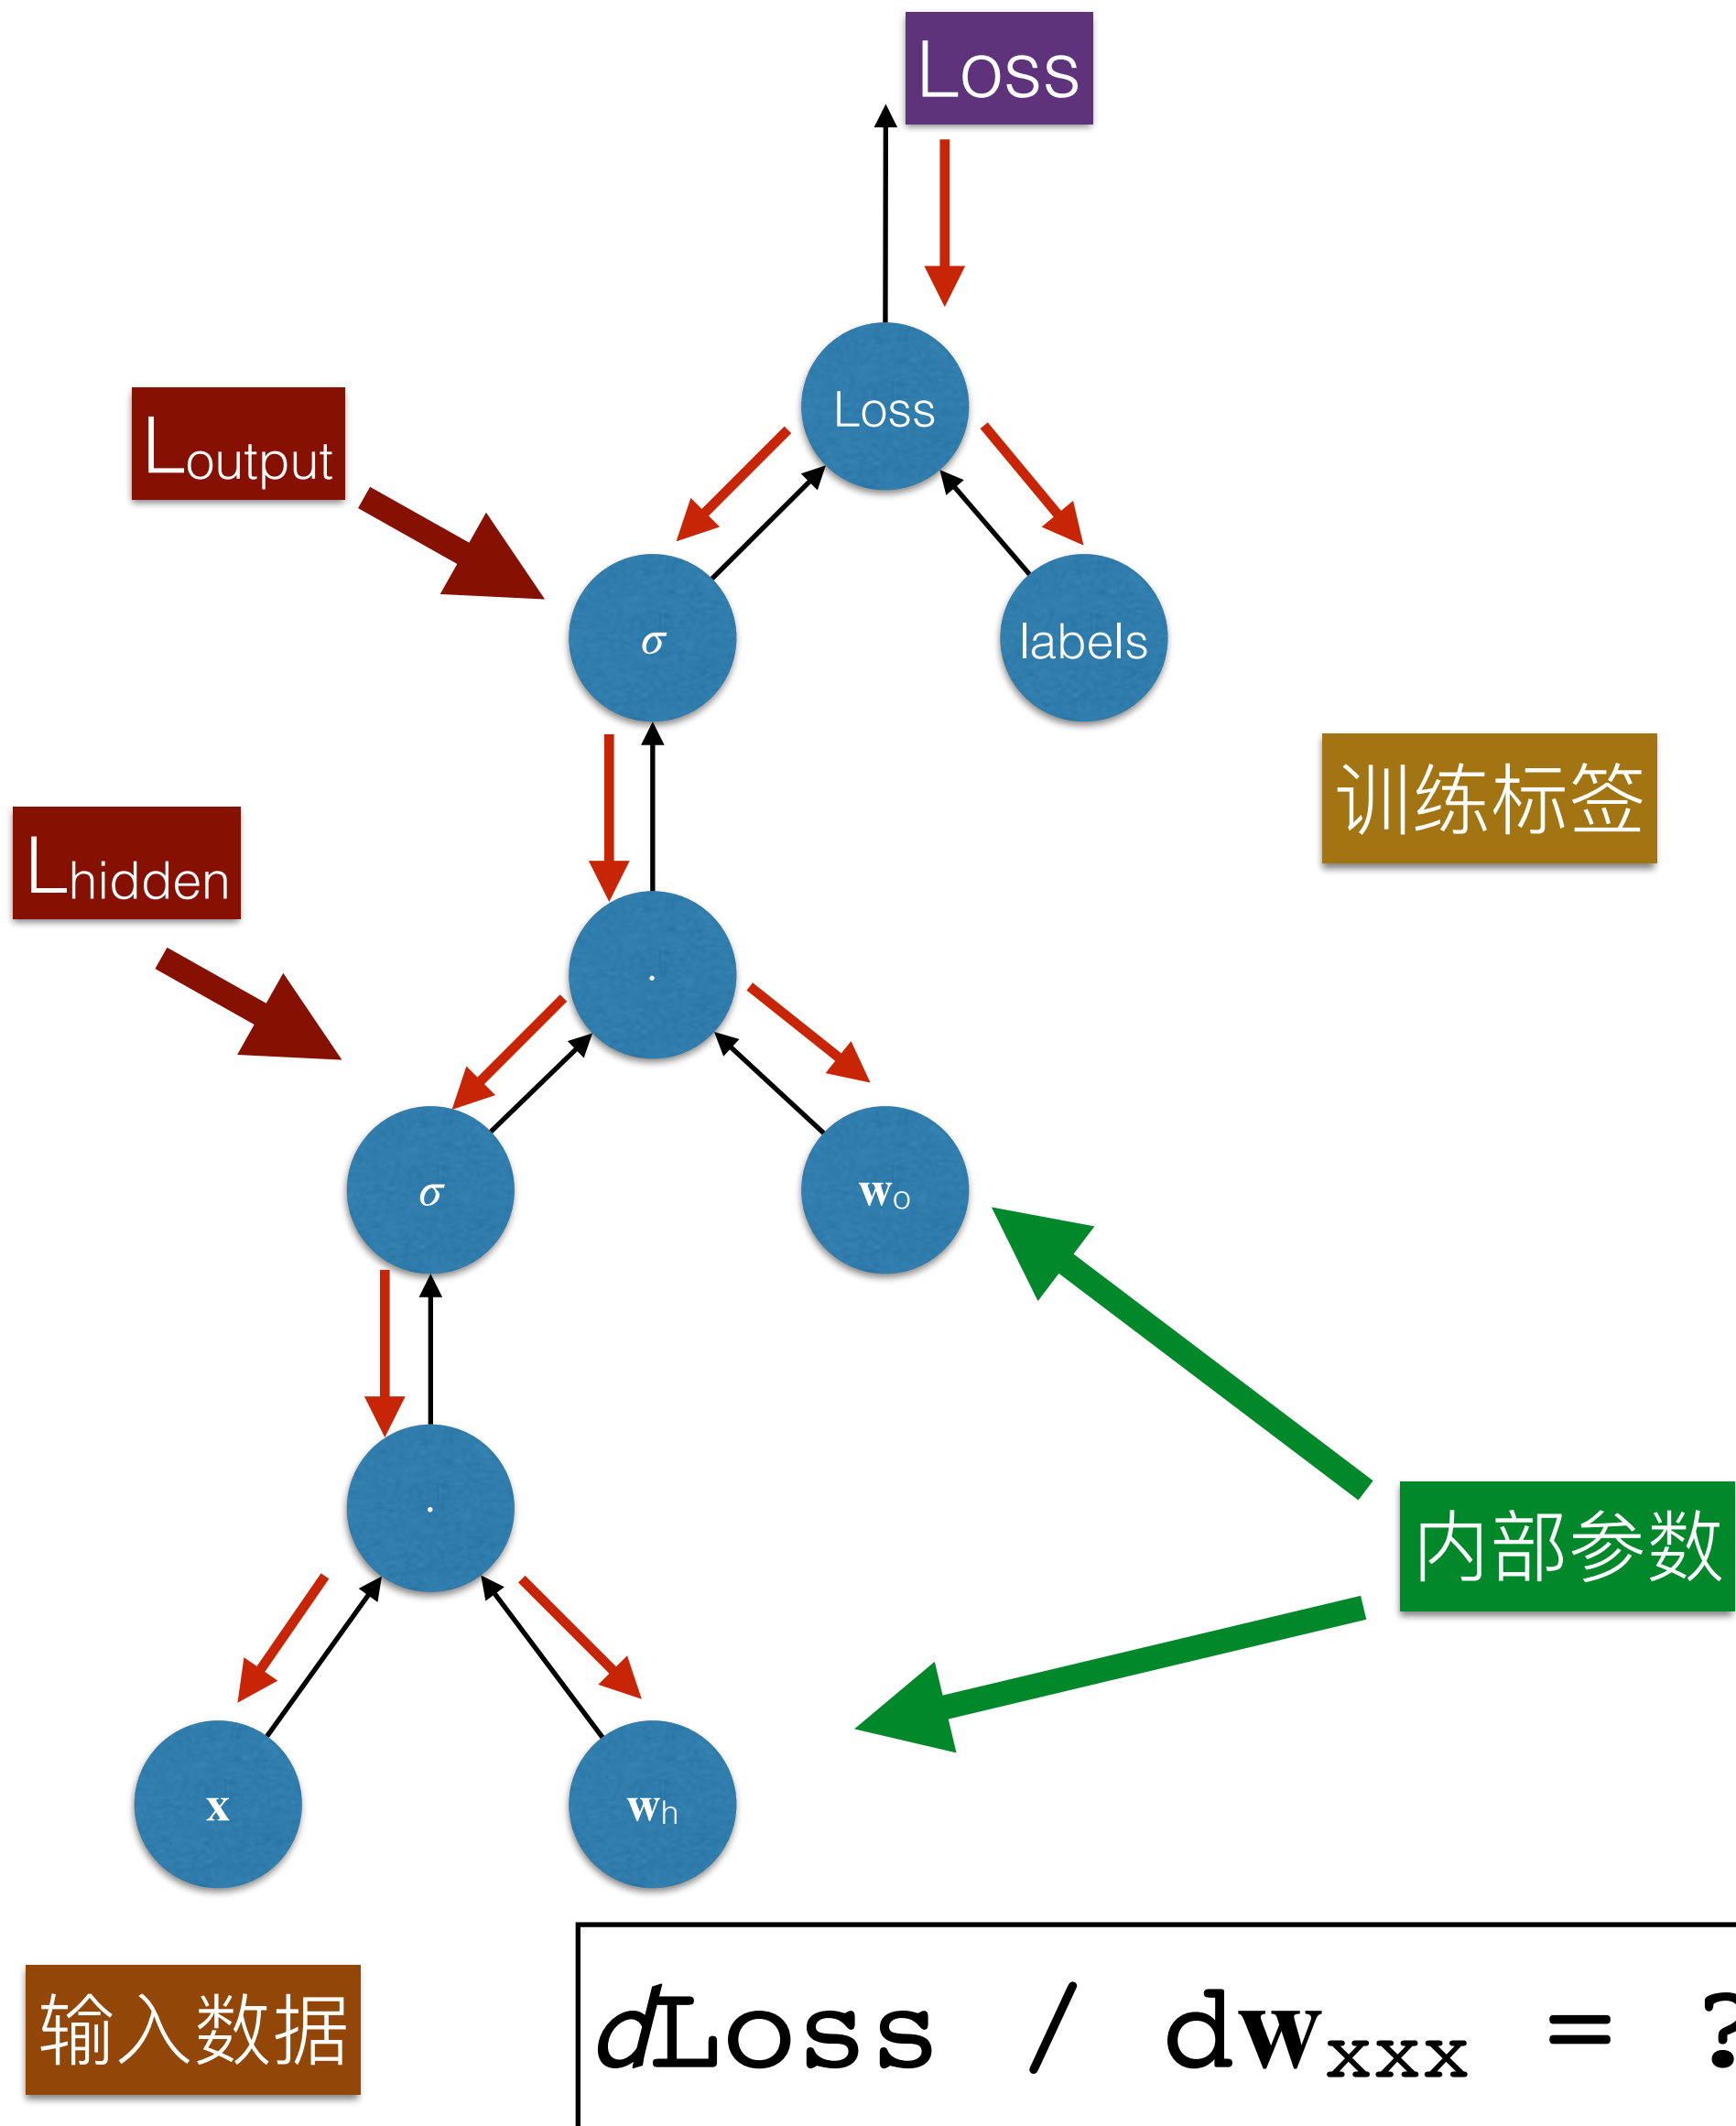
$$L_{\text{output}} = x \cdot W_{\text{output}} + b_{\text{output}}$$

$$\text{Loss} = \text{cross_entropy}(L_{\text{output}}, \text{labels})$$

通过大量的 x 和labels, 找到一组 $W_{\text{hidden}}, W_{\text{output}}, b_{\text{hidden}}, b_{\text{output}}$, 使得Loss最小。

反向传播

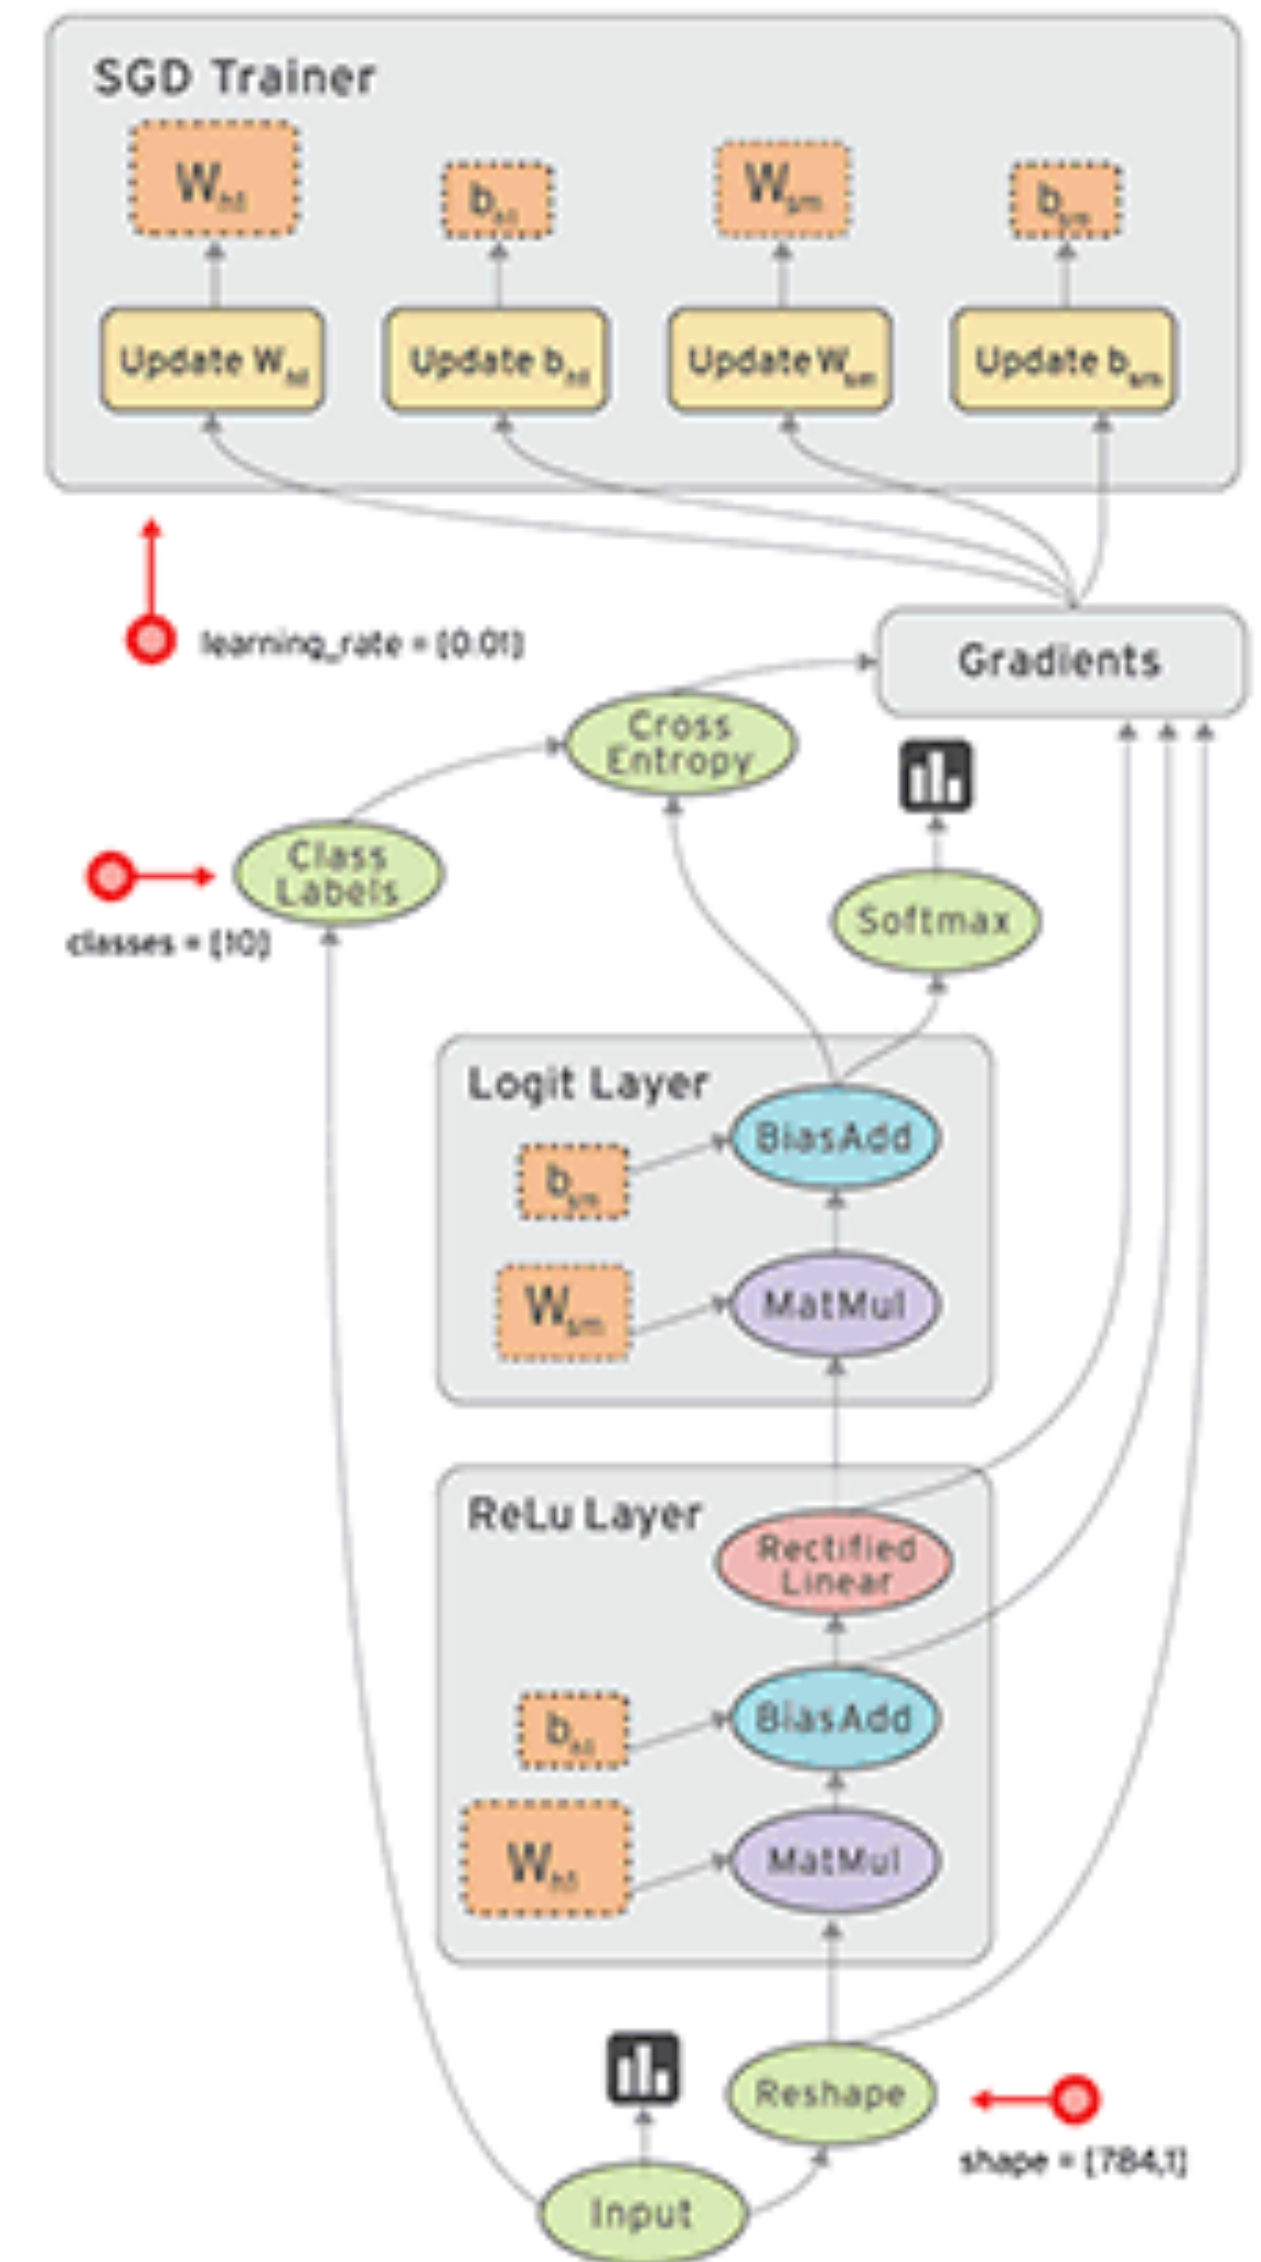
$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial x}$$



$$d\text{Loss} / dw_o =$$
$$(d\text{Loss} / dL_{\text{output}}) \cdot (dL_{\text{output}} / dw_o)$$

$$d\text{Loss} / dw_h =$$
$$(d\text{Loss} / dL_{\text{output}}) \cdot (dL_{\text{output}} / dL_{\text{hidden}}) \cdot$$
$$(dL_{\text{hidden}} / dw_h)$$

- Google Brain开发维护
 - Jeff Dean (BigTable, MapReduce)
 - Yangqing Jia (Caffe)
 - Ian Goodfellow (Theano)
- 声明式定义计算图
- 自动求导
- 支持集群计算
 - 模型并行
 - 数据并行
- 工程支持
 - TensorBoard
 - TensorFlow Serving
 - Integrated with Kubernetes



训练数据

```
import tensorflow as tf
```

```
X = tf.constant([[0, 0, 1], [0, 1, 1], [1, 0, 1], [1, 1, 1]], dtype=tf.float32)
y = tf.constant([0, 1, 1, 0])
```

```
syn0 = tf.Variable(tf.random_uniform([3, 4], minval=-1, maxval=1, dtype=tf.float32))
syn1 = tf.Variable(tf.random_uniform([4, 2], minval=-1, maxval=1, dtype=tf.float32))
```

```
l1 = tf.sigmoid(tf.matmul(X, syn0))
l2 = tf.sigmoid(tf.matmul(l1, syn1))
```

声明计算逻辑

声明参数

```
loss_op = tf.reduce_mean(tf.nn.sparse_softmax_cross_entropy_with_logits(l2, y))
train_op = tf.train.GradientDescentOptimizer(0.1).minimize(loss_op)
```

```
init = tf.initialize_all_variables()
sess = tf.Session()
sess.run(init)
```

声明反向计算逻辑和参数更新策略

```
for i in range(60000):
    _, loss = sess.run([train_op, loss_op])
```

执行计算



练习1 - 线性回归



练习环境

- <https://github.com/tongda/tf-tutorial>
- Python
- TensorFlow
- ipython notebook
- matplotlib



练习2 - 神经网络

多层感知机——任务定义



$$L_{\text{hidden}} = \text{sigmoid}(x \cdot W_{\text{hidden}} + b_{\text{hidden}})$$

$$L_{\text{output}} = x \cdot W_{\text{output}} + b_{\text{output}}$$

$$\text{Loss} = \text{cross_entropy}(L_{\text{output}}, \text{labels})$$



通过大量的 x 和labels, 找到一组 $W_{\text{hidden}}, W_{\text{output}}, b_{\text{hidden}}, b_{\text{output}}$, 使得Loss最小。

激活函数

Activation Functions


Sigmoid

$$\sigma(x) = 1 / (1 + e^{-x})$$




tanh tanh(x)

ReLU max(0,x)



Leaky ReLU max(0.1x, x)




Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$f(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha (\exp(x) - 1) & \text{if } x \leq 0 \end{cases}$$





激活函数

Activation Functions


Sigmoid

$$\sigma(x) = 1 / (1 + e^{-x})$$




tanh $\tanh(x)$

ReLU $\max(0, x)$



Leaky ReLU

$$\max(0.1x, x)$$





Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$f(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha (\exp(x) - 1) & \text{if } x \leq 0 \end{cases}$$






卷积神经网络


Convolution Layer

32x32x3 image




Convolution Layer

32x32x3 image




5x5x3 filter



用滤波器/滤镜/过滤器 (filter) 对图片执行卷积操作。卷积操作是将滤波器在图片上进行滑动，每滑动一个位置，计算滤波器和图片上对应位置多维矩阵的点积。


Convolution Layer

32x32x3 image




卷积核的深度（即图片的通道数目），和输入图片的深度保持一致。

5x5x3 filter



用滤波器/滤镜/过滤器/卷积核（filter）对图片执行卷积操作。卷积操作是将卷积核在图片上进行滑动，每滑动一个位置，计算滤波器和图片上对应位置多维矩阵的点积。

Convolution Layer



Convolution Layer




Convolution Layer

再添加另一个卷积核，见图中的绿色部分。




假如我们有6个 $5 \times 5 \times 3$ 的卷积核，就会得到6个激活地图（activation map）。




将6个激活地图叠加在一起，就得到一个新的图片，尺寸是 $28 \times 28 \times 6$ 。

卷积网络 (ConvNet) 就是一系列卷积层，通过激活函数连接在一起。

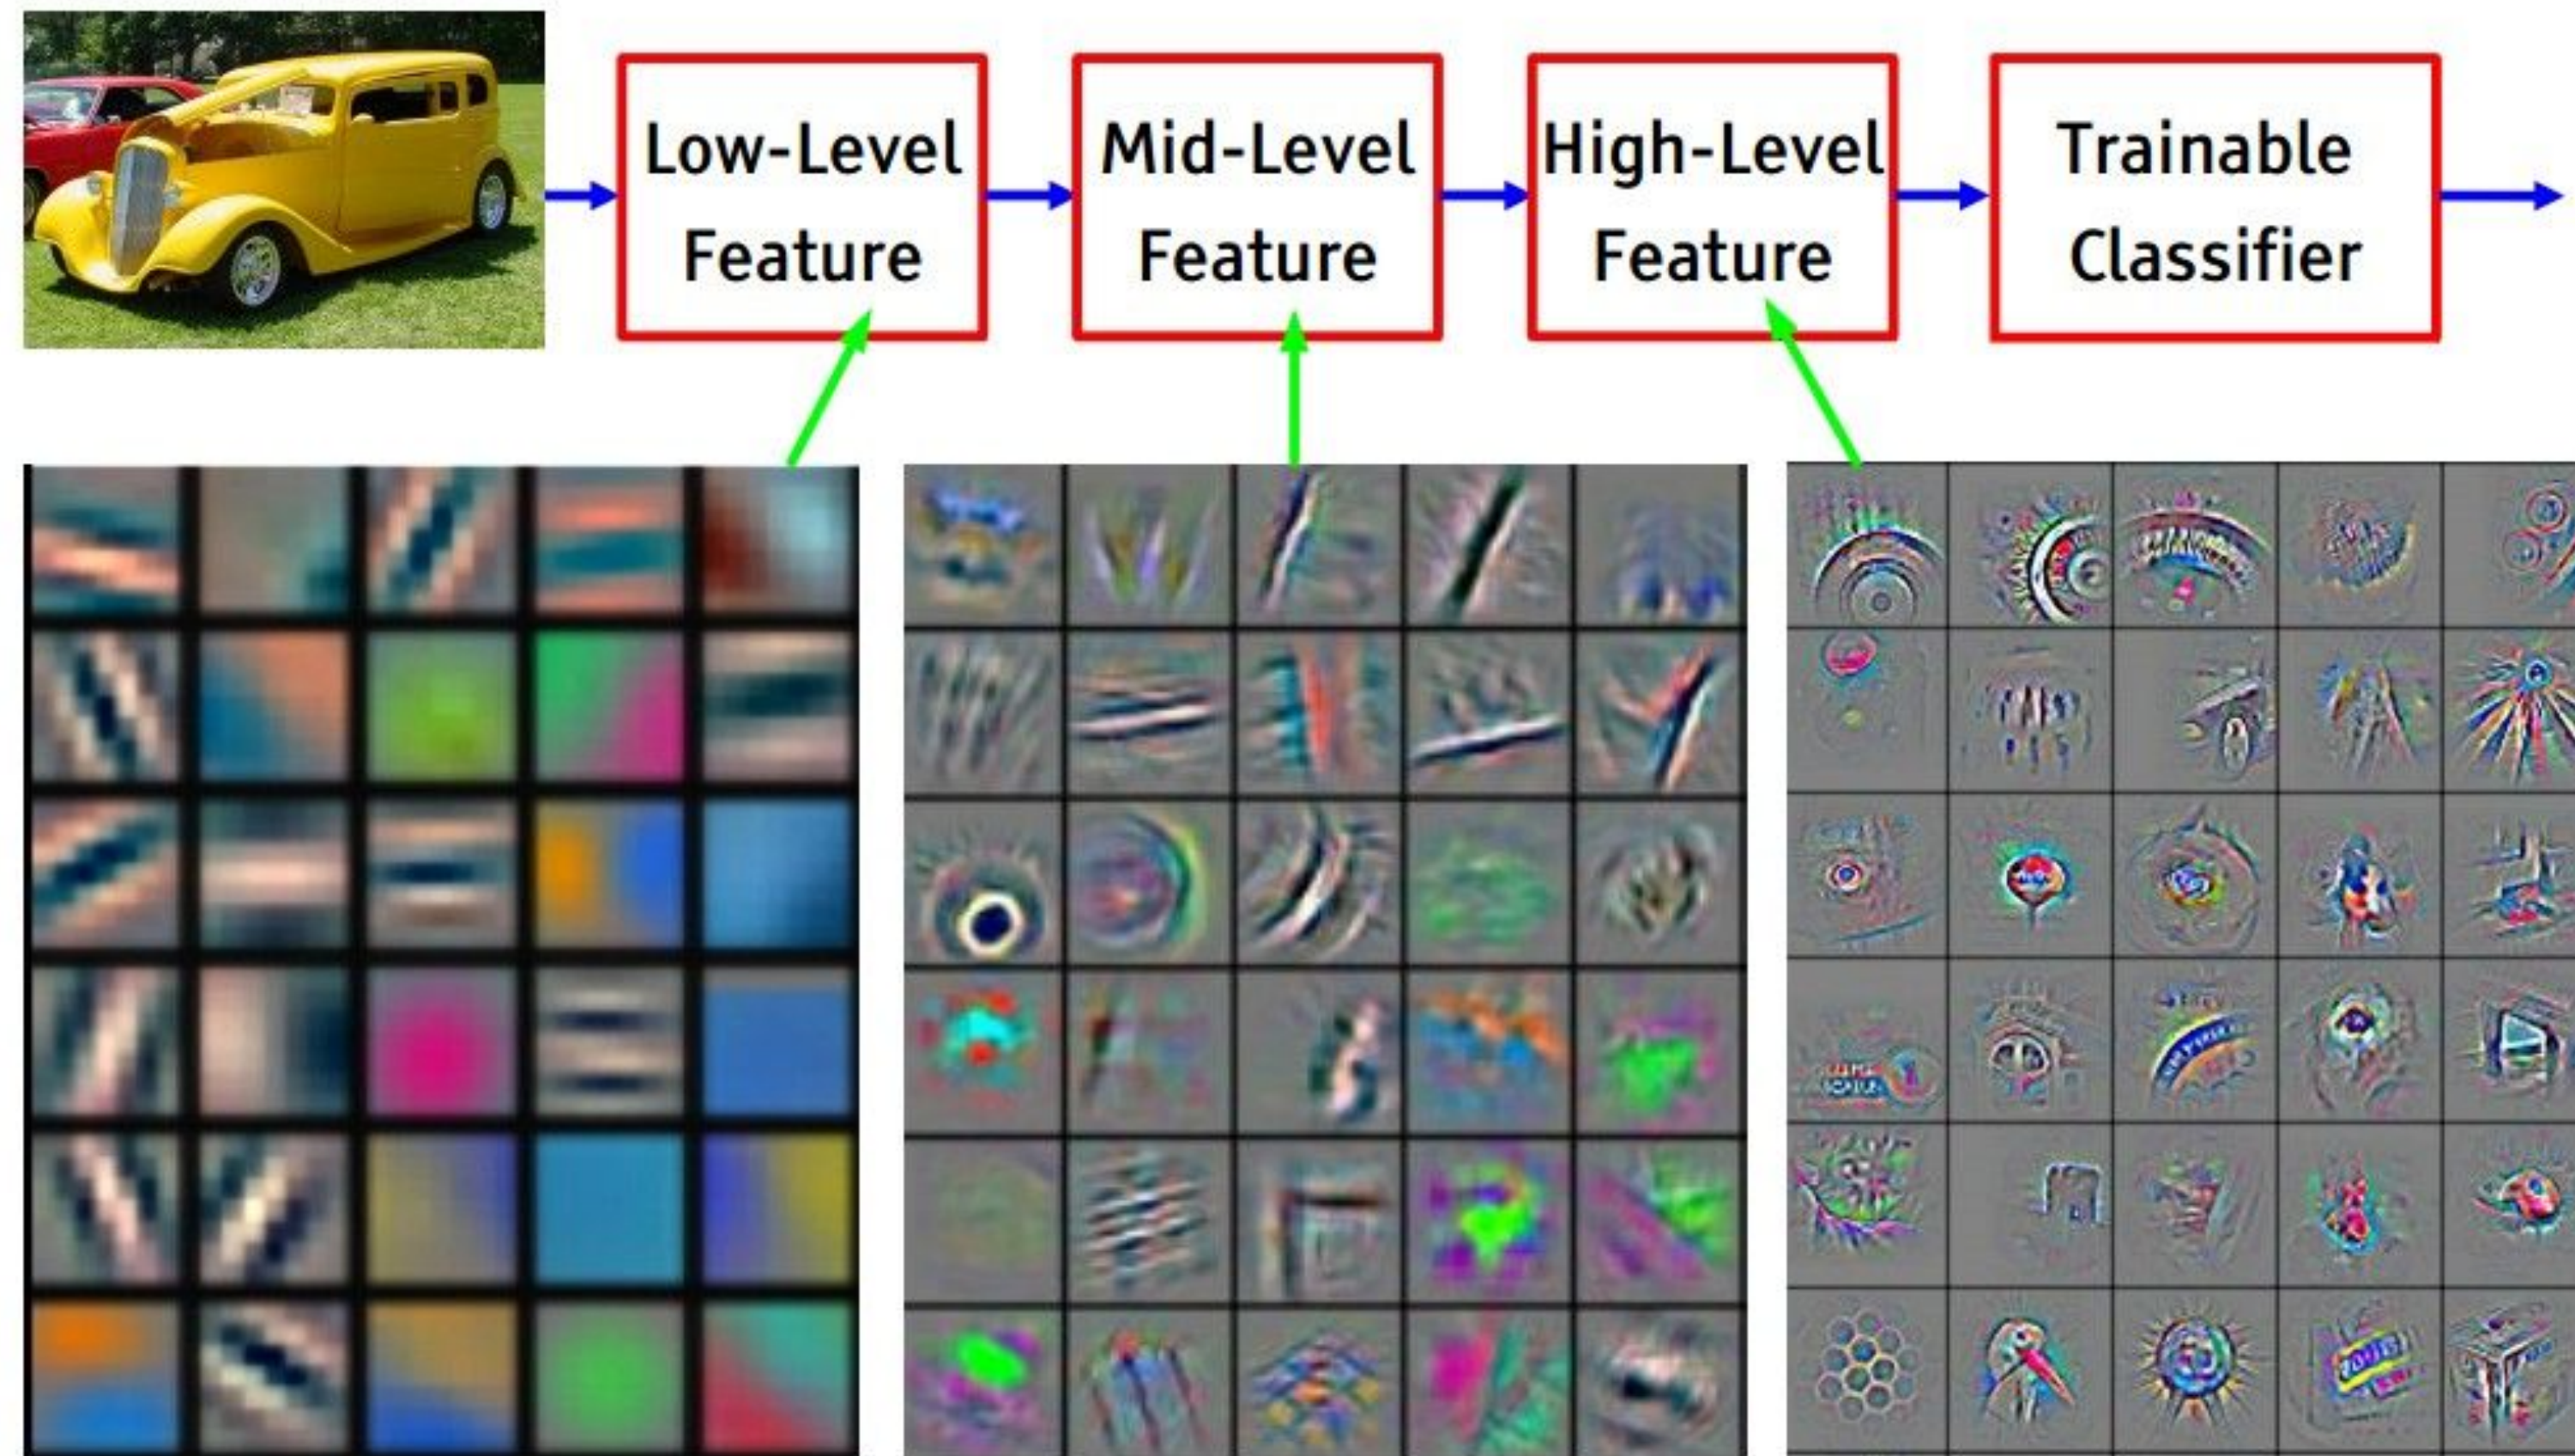


卷积网络 (ConvNet) 就是一系列卷积层，通过激活函数连接在一起。



Preview

[From recent Yann LeCun slides]



Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]


Stride & Padding

仔细看一下空间维度的变化。



Stride & Padding


仔细看一下空间维度的变化。



7x7 input (spatially)
assume 3x3 filter

Stride & Padding


仔细看一下空间维度的变化。



7x7 input (spatially)
assume 3x3 filter

Stride & Padding


仔细看一下空间维度的变化。



7x7 input (spatially)
assume 3x3 filter

Stride & Padding


仔细看一下空间维度的变化。



7x7 input (spatially)
assume 3x3 filter

Stride & Padding

仔细看一下空间维度的变化。



7x7 input (spatially)
assume 3x3 filter

=> 5x5 output

Stride & Padding

实际使用中，通常用0来填充边缘。

0	0	0	0	0	0			
0								
0								
0								
0								

e.g. input 7x7

3x3 filter, applied with **stride 1**

pad with 1 pixel border => what is the output?


7x7 output!

常见的做法是步进为1的卷积层，卷积核为 $F \times F$ ，填充的0就是 $(F-1)/2$ ，这样可以保证输出的图片大小不变。

e.g. $F = 3 \Rightarrow$ zero pad with 1

$F = 5 \Rightarrow$ zero pad with 2

$F = 7 \Rightarrow$ zero pad with 3




卷积神经网络 - 续

池化

Pooling layer


- makes the representations smaller and more manageable
- operates over each activation map independently:




池化

MAX POOLING

Single depth slice



max pool with 2x2 filters
and stride 2

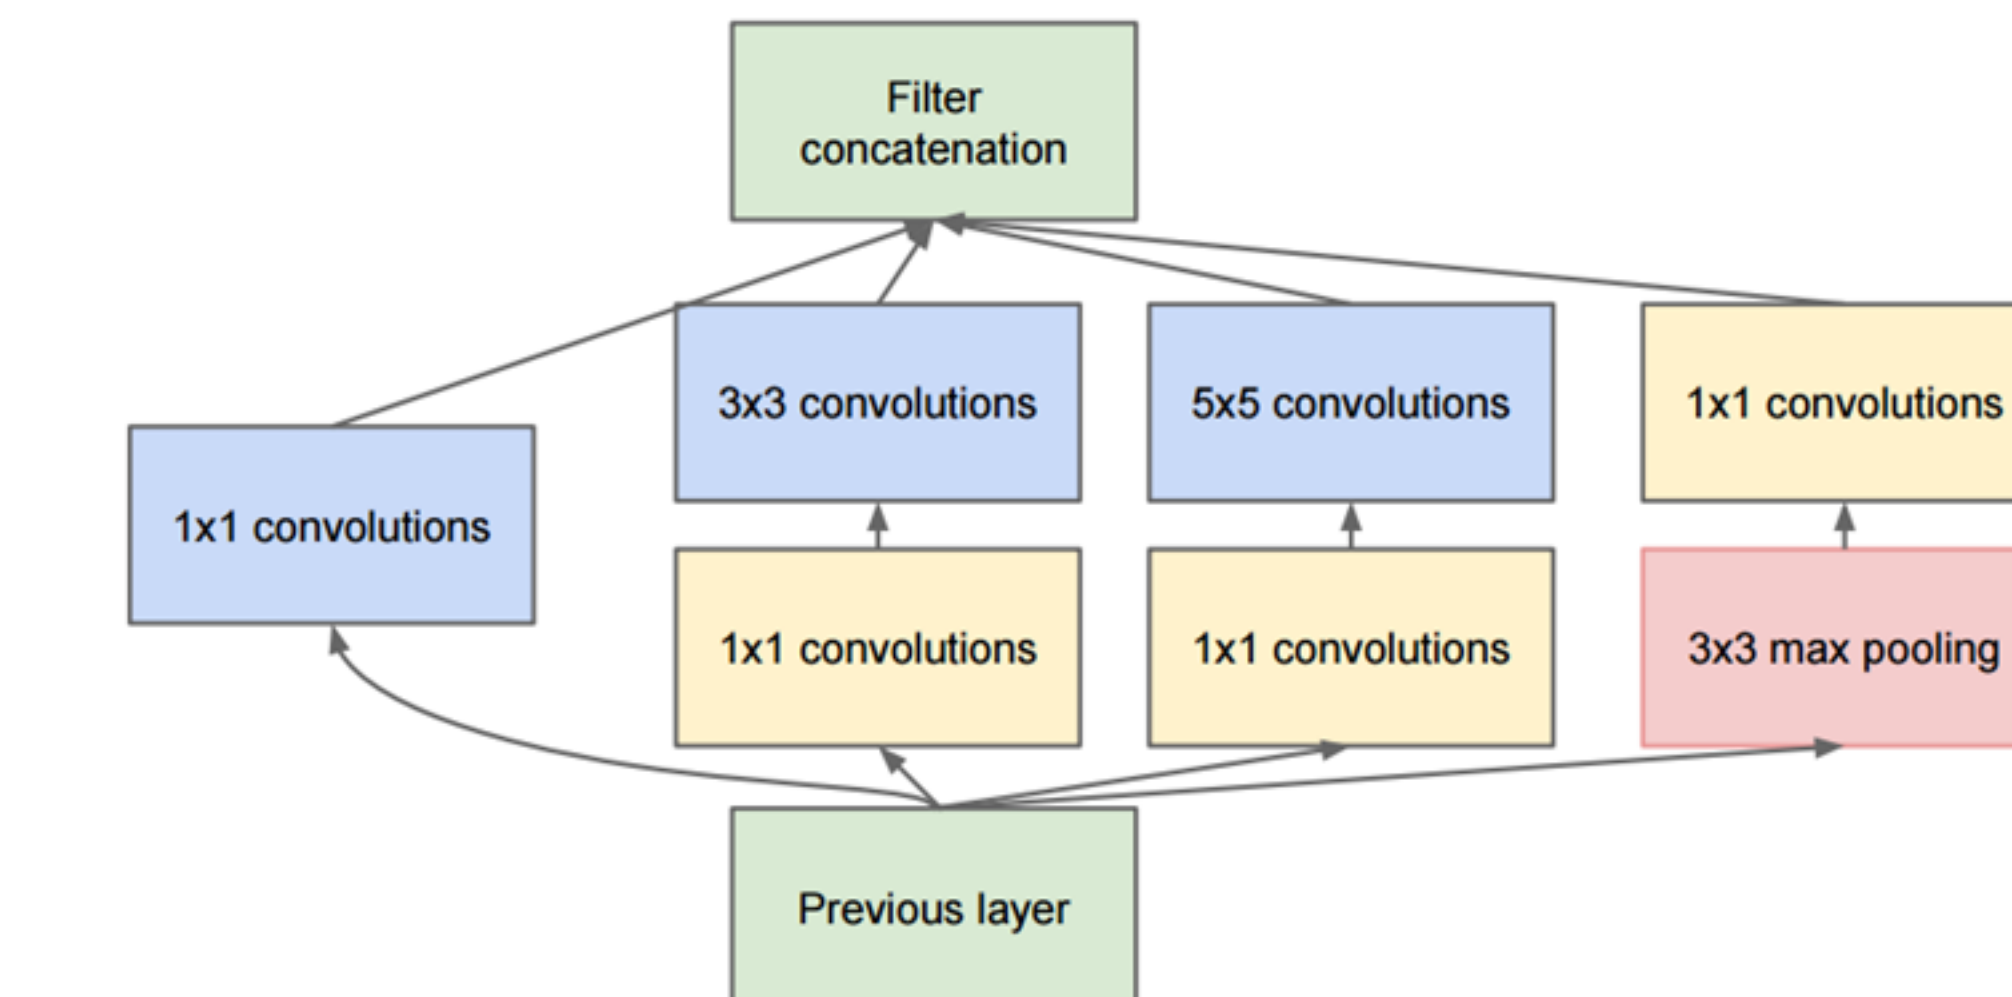



Regularization


- L2 Regularization
- Drop out
- Batch Norm

其他结构

- Inception
- ResNet




Full Inception module



A residual block

论文中的卷积网络

two more layers to go: POOL/FC



论文中的卷积网络

Case Study: AlexNet

[Krizhevsky et al. 2012]

Full (simplified) AlexNet architecture:

[227x227x3] INPUT

[55x55x96] **CONV1**: 96 11x11 filters at stride 4, pad 0

[27x27x96] **MAX POOL1**: 3x3 filters at stride 2

[27x27x96] **NORM1**: Normalization layer

[27x27x256] **CONV2**: 256 5x5 filters at stride 1, pad 2

[13x13x256] **MAX POOL2**: 3x3 filters at stride 2

[13x13x256] **NORM2**: Normalization layer

[13x13x384] **CONV3**: 384 3x3 filters at stride 1, pad 1

[13x13x384] **CONV4**: 384 3x3 filters at stride 1, pad 1


[13x13x256] **CONV5**: 256 3x3 filters at stride 1, pad 1

[6x6x256] **MAX POOL3**: 3x3 filters at stride 2

[4096] **FC6**: 4096 neurons

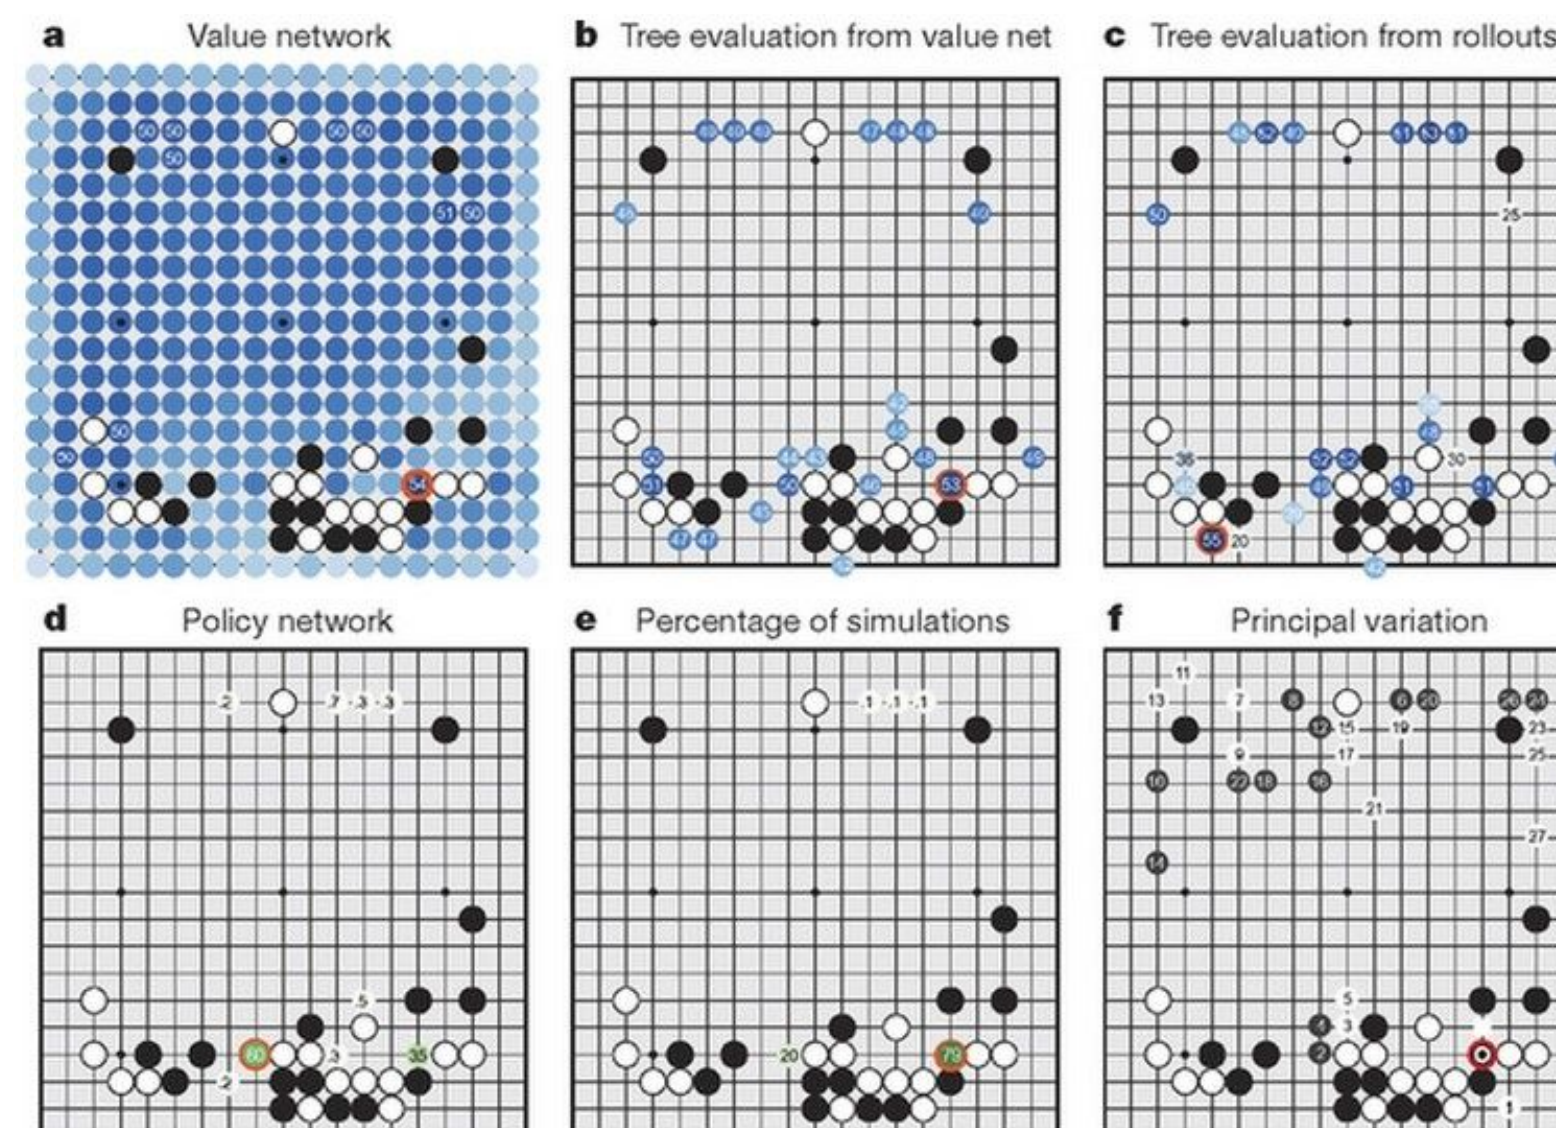
[4096] **FC7**: 4096 neurons

[1000] **FC8**: 1000 neurons (class scores)



论文中的卷积网络

Case Study Bonus: DeepMind's AlphaGo



论文中的卷积网络

The input to the policy network is a $19 \times 19 \times 48$ image stack consisting of 48 feature planes. The first hidden layer zero pads the input into a 23×23 image, then convolves k filters of kernel size 5×5 with stride 1 with the input image and applies a rectifier nonlinearity. Each of the subsequent hidden layers 2 to 12 zero pads the respective previous hidden layer into a 21×21 image, then convolves k filters of kernel size 3×3 with stride 1, again followed by a rectifier nonlinearity. The final layer convolves 1 filter of kernel size 1×1 with stride 1, with a different bias for each position, and applies a softmax function. The match version of AlphaGo used $k = 192$ filters; [Fig. 2b](#) and [Extended Data Table 3](#) additionally show the results of training with $k = 128, 256$ and 384 filters.

policy network:

[19x19x48] Input

CONV1: 192 5x5 filters , stride 1, pad 2 => [19x19x192]

CONV2..12: 192 3x3 filters, stride 1, pad 1 => [19x19x192]

CONV: 1 1x1 filter, stride 1, pad 0 => [19x19] (*probability map of promising moves*)



练习

对比不同模型性能

- Model Evaluation
- Cross Validation

对比不同模型性能

```
# Build graph to count correct answers.
def mnist_evaluation(logits, labels):
    """Evaluate the quality of the logits at predicting the label.
    Args:
        logits: Logits tensor, float - [BATCH_SIZE, NUM_CLASSES].
        labels: Labels tensor, int32 - [BATCH_SIZE], with values in the
        range [0, NUM_CLASSES).
    Returns:
        A scalar int32 tensor with the number of examples (out of batch_size)
        that were predicted correctly.
    """
    # For a classifier model, we can use the in_top_k Op.
    # It returns a bool tensor with shape [batch_size] that is true for
    # the examples where the label is in the top k (here k=1)
    # of all logits for that example.
    correct = tf.nn.in_top_k(logits, labels, 1)
    # Return the number of true entries.
    return tf.reduce_sum(tf.cast(correct, tf.int32))
```

对比不同模型性能

```
# Print out loss value.
if step % 100 == 0:
    eval_value = sess.run(eval_op,
                           feed_dict={images_placeholder: images_feed,
                                       labels_placeholder: labels_feed})
    print('Step %d: loss = %.2f, eval = %.2f' % (step, loss_value, eval_value / BATCH_SIZE))
```